Stat 6021: HW 1

Tom Lever

08/26/22

Download the dataset `UScovid.csv` from Collab. The dataset was released by *The New York Times* and contains data on cumulative (i.e., accruing) counts of coronavirus cases and deaths in the United States, at the state and county level, over each day from Jan 21, 2020 to June 3, 2021. You may read more about the data and the variable descriptions here Please note the dataset is regularly updated. We will use the file on Collab.

Read the data file into R and store the dataset into the object Covid [Text].

```
Covid <- read.csv("USCovid.csv")
head(Covid, n = 3)
```

```
##         date    county      state  fips cases deaths
## 1 2020-01-21 Snohomish Washington 53061     1      0
## 2 2020-01-22 Snohomish Washington 53061     1      0
## 3 2020-01-23 Snohomish Washington 53061     1      0
```

```
nrow(Covid)
```

```
## [1] 1384683
```

There are 1,384,683 snapshots in this dataset. The header row of `Covid` is not considered in this determination.

1. For this question, we focus on data at the county level.

   (a) We are interested in the data at the most recent date, June 3, 2021 (i.e., 2021-06-03). Create a data frame called `latest` that

   - has only rows pertaining to data from June 3, 2021,
   - removes rows pertaining to counties that are "Unknown",
   - removes the column `date` and `fips`, and
   - is ordered by `county` and then `state` alphabetically.

   Use the `head()` function to display the first 6 rows of the data frame `latest`.

   ```
   library(dplyr)
   latest <-
       Covid%>%
           filter(date == "2021-06-03")%>%
           filter(county != "Unknown")%>%
           # filter(!is.na(county))
           # Unknown is not equivalent to NA (i.e., Not Available)
           select(-date, -fips)%>%
           arrange(county, state)
   head(latest, n = 6)
   ```

   ```
   ##      county          state cases deaths
   ## 1 Abbeville South Carolina  2599     41
   ## 2    Acadia      Louisiana  6703    195
   ```

```
## 3  Accomack       Virginia  2862     43
## 4       Ada          Idaho 52964    475
## 5     Adair           Iowa   873     32
## 6     Adair       Kentucky  1944     54
```

(b) Calculate the death rate– call it `death.rate` –for each county. Report the death rate as a percent and round to two decimal places. Add `death.rate` as a new column to the data frame `latest`. Display the first 6 rows of the data frame `latest`.

```
death.rate <- round(latest%>%select(deaths) / latest%>%select(cases) * 100, 2)
colnames(death.rate) <- "death.rate"
latest <- bind_cols(latest, death.rate)

death.rate <-
    rename(
        round(latest%>%select(deaths) / latest%>%select(cases) * 100, 2),
        death.rate = deaths
    )
latest <- latest%>%mutate(death.rate = death.rate)
head(latest, n = 6)
```

```
##      county          state cases deaths death.rate
## 1 Abbeville South Carolina  2599     41       1.58
## 2    Acadia      Louisiana  6703    195       2.91
## 3  Accomack       Virginia  2862     43       1.50
## 4       Ada          Idaho 52964    475       0.90
## 5     Adair           Iowa   873     32       3.67
## 6     Adair       Kentucky  1944     54       2.78
```

(c) Display the counties with the 10 largest numbers of cases. Be sure to display also the appropriate states, numbers of deaths, and death rates.

```
# "slice_min() and slice_max() can order_by multiple variables
# if you supply them as a data.frame or tibble (#6176)."
# https://github.com/tidyverse/dplyr/blob/main/NEWS.md
# For slice_max, devtools::install_github("tidyverse/dplyr")
slice_max(
    latest,
    n = 10,
    order_by = data.frame(cases, county, state, deaths, death.rate),
    with_ties = FALSE
)
```

```
##            county      state   cases deaths death.rate
## 1     Los Angeles California 1245127  24375       1.96
## 2   New York City   New York  949986  33257       3.50
## 3            Cook   Illinois  554390  10893       1.96
## 4        Maricopa    Arizona  551509  10084       1.83
## 5      Miami-Dade    Florida  501925   6472       1.29
## 6          Harris      Texas  401345   6462       1.61
## 7          Dallas      Texas  303533   4082       1.34
## 8       Riverside California  300879   4614       1.53
## 9  San Bernardino California  298599   4760       1.59
## 10      San Diego California  280410   3760       1.34
```

(d) Display the counties with the 10 largest numbers of deaths. Be sure to display also the appropriate states, numbers of cases, and death rates.

```
slice_max(
    latest,
    n = 10,
    order_by = data.frame(deaths, county, state, cases, death.rate),
    with_ties = FALSE
)
```

```
##            county      state   cases deaths death.rate
## 1   New York City   New York  949986  33257       3.50
## 2     Los Angeles California 1245127  24375       1.96
## 3            Cook   Illinois  554390  10893       1.96
## 4        Maricopa    Arizona  551509  10084       1.83
## 5      Miami-Dade    Florida  501925   6472       1.29
## 6          Harris      Texas  401345   6462       1.61
## 7          Orange California  272242   5070       1.86
## 8           Wayne   Michigan  164612   5048       3.07
## 9  San Bernardino California  298599   4760       1.59
## 10      Riverside California  300879   4614       1.53
```

(e) Display the counties with the 10 highest death rates. Be sure to display also the appropriate states, numbers of cases, and numbers of deaths. Is there something you notice about these counties?

```
# For calculate_percentile,
# devtools::install_github("tslever/Tom_Levers_Git_Repository/TomLeversRPackage")
library(TomLeversRPackage)
counties_with_10_highest_death_rates <-
    slice_max(
        latest,
        n = 10,
        order_by = data.frame(death.rate, county, state, cases, deaths),
        with_ties = FALSE
    )
counties_with_10_highest_death_rates
```

```
##          county      state cases deaths death.rate
## 1         Grant   Nebraska    41      4       9.76
## 2        Sabine      Texas   524     45       8.59
## 3     Petroleum    Montana    12      1       8.33
## 4       Harding New Mexico    12      1       8.33
## 5         Foard      Texas   124     10       8.06
## 6       Hancock    Georgia   928     68       7.33
## 7      Glascock    Georgia   269     19       7.06
## 8        Motley      Texas   116      8       6.90
## 9  Throckmorton      Texas    73      5       6.85
## 10      Candler    Georgia   978     67       6.85
```

```
calculate_percentile(
    latest%>%pull(cases),
    max(counties_with_10_highest_death_rates%>%select(cases), na.rm = TRUE)
)
```

```
## [1] 23
```

Yes. The percentile of the maximum number of cases among the counties with the 10 highest death rates, given all numbers of cases, is 23. The maximum number of cases among the counties with the 10 highest death rates is in the lowest quarter of numbers of cases.

(f) Display the counties with the 10 highest death rates among counties with at least 100,000 cases. Be sure to display also the appropriate states, numbers of cases, and numbers of deaths.

```
library(TomLeversRPackage)
latest%>%
    filter(cases > 100000)%>%
    slice_max(
        n = 10,
        order_by = data.frame(death.rate, county, state, cases, deaths),
        with_ties = FALSE
    )
```

```
##           county         state  cases deaths death.rate
## 1  New York City      New York 949986  33257       3.50
## 2          Wayne      Michigan 164612   5048       3.07
## 3      Middlesex Massachusetts 134980   3761       2.79
## 4         Bergen    New Jersey 104301   2868       2.75
## 5         Macomb      Michigan 100190   2441       2.44
## 6   Philadelphia  Pennsylvania 153521   3692       2.40
## 7      St. Louis      Missouri 100195   2249       2.24
## 8      Fairfield   Connecticut 100093   2198       2.20
## 9           Pima       Arizona 116997   2406       2.06
## 10       Oakland      Michigan 118035   2368       2.01
```

(g) Display the number of cases, deaths, and death rate for the following counties.

i. Albemarle, Virginia

```
latest%>%
    filter(county == "Albemarle" & state == "Virginia")
```

```
##      county    state cases deaths death.rate
## 1 Albemarle Virginia  5801     83       1.43
```

ii. Charlottesville City, Virginia

```
latest%>%
    filter(county == "Charlottesville city" & state == "Virginia")
```

```
##                 county    state cases deaths death.rate
## 1 Charlottesville city Virginia  4014     57       1.42
```

2. For this question, we focus on data at the state level. Note that the dataset has data on the 50 states, plus DC, Puerto Rico, Guam, Northern Mariana Islands, and the Virgin Islands.

(a) We are interested in the data at the most recent date, June 3, 2021. Create a data frame called state.level that

- has 55 rows, including 1 for each state, 1 for DC, and 1 for each territory
- has 3 columns, including `state`, `cases`, and `deaths`, and
- is ordered alphabetically by `state`.

Display the first 6 rows of the data frame `state.level`.

```
state.level <-
    Covid%>%
        filter(date == "2021-06-03")%>%
        group_by(state)%>%
        summarize(cases = sum(cases), deaths = sum(deaths, na.rm = TRUE))
head(state.level, n = 6)
```

```
## # A tibble: 6 x 3
##   state        cases deaths
##   <chr>        <int>  <int>
## 1 Alabama     545028  11188
## 2 Alaska       69826    352
## 3 Arizona     882691  17653
## 4 Arkansas    341889   5842
## 5 California 3793055  63345
## 6 Colorado    547961   6746
```

```
nrow(state.level)
```

```
## [1] 55
```

(b) Calculate the death rate (call it `state.rate`). Report the death rate as a percent and round to two decimal places. Add `state.rate` as a new column to the data frame `state.level`. Display the first 6 rows of the data frame `state.level`.

```
state.rate <-
    round(state.level%>%select(deaths) / state.level%>%select(cases) * 100, 2)
colnames(state.rate) <- "state.rate"
state.level <- bind_cols(state.level, state.rate)

state.rate <-
    rename(
        round(
            state.level%>%select(deaths) / state.level%>%select(cases) * 100,
            2
        ),
        state.rate = deaths
    )
state.level <- state.level%>%mutate(state.rate = state.rate)
head(state.level, n = 6)
```

```
## # A tibble: 6 x 4
##   state        cases deaths state.rate
##   <chr>        <int>  <int>      <dbl>
## 1 Alabama     545028  11188       2.05
## 2 Alaska       69826    352       0.5
## 3 Arizona     882691  17653       2
## 4 Arkansas    341889   5842       1.71
## 5 California 3793055  63345       1.67
## 6 Colorado    547961   6746       1.23
```

(c) What is the death rate in Virginia?

```
state.level%>%
    filter(state == "Virginia")%>%
    select(state, state.rate)
```

```
## # A tibble: 1 x 2
##   state    state.rate
##   <chr>         <dbl>
## 1 Virginia       1.66
```

The death rate in Virginia is 1.66 percent.

(d) What is the death rate in Puerto Rico?

```
state.level%>%
    filter(state == "Puerto Rico")%>%
    select(state, state.rate)
```

```
## # A tibble: 1 x 2
##   state       state.rate
##   <chr>            <dbl>
## 1 Puerto Rico       1.46
```

The death rate in Puerto Rico is 1.46 percent.

(e) Which states have the 10 highest death rates?

```
slice_max(
    state.level,
    n = 10,
    order_by = data.frame(state.rate, state, cases, deaths),
    with_ties = FALSE
)
```

```
## # A tibble: 10 x 4
##    state                  cases deaths state.rate
##    <chr>                  <int>  <int>      <dbl>
##  1 New Jersey           1017044  26253       2.58
##  2 Massachusetts         707523  17893       2.53
##  3 New York             2102003  52811       2.51
##  4 Connecticut           347748   8245       2.37
##  5 District of Columbia   49041   1136       2.32
##  6 Mississippi           318048   7324       2.3
##  7 Pennsylvania         1208879  27349       2.26
##  8 Louisiana             472617  10605       2.24
##  9 New Mexico            203330   4275       2.1
## 10 Maryland              460406   9626       2.09
```

The states with the 10 highest death rates are listed in the above column `state`.

(f) Which states have the 10 lowest death rates?

```
slice_min(
    state.level,
    n = 10,
    order_by = data.frame(state.rate, state, cases, deaths),
    with_ties = FALSE
)
```

```
## # A tibble: 10 x 4
##    state                     cases deaths state.rate
##    <chr>                     <int>  <int>      <dbl>
##  1 Alaska                    69826    352       0.5
##  2 Utah                     406895   2308       0.57
##  3 Virgin Islands             3512     28       0.8
##  4 Vermont                   24240    255       1.05
##  5 Nebraska                 223517   2385       1.07
##  6 Idaho                    192704   2103       1.09
##  7 Northern Mariana Islands    183      2       1.09
##  8 Wisconsin                675152   7923       1.17
##  9 Wyoming                   60543    720       1.19
```

```
## 10 Colorado                  547961   6746       1.23
```

The states with the 10 lowest death rates are listed in the above column `state`.

(g) Export this dataset as a `.csv` file named `stateCovid.csv`. We will be using this file for the next homework.

I assume "this dataset" is `state.level`.

```
write.csv(state.level, "stateCovid.csv", row.names = FALSE)
```